All-hands meeting notes

Planning call for the UvA researchers' AI course, Thursday 8 October 2026

Who	Gjorgji and Joaquim, plus at least two more from the course team. The transcript has four anonymous speakers, so the rest is inferred. Joaquim left partway through.
Source	Recording and transcript in this folder, plus the auto-generated summary.
Status	Notes written from the transcript. Items marked *to confirm* were garbled or ambiguous in it.

Summary

The team settled how participants get their software and licenses, and walked through the content of each evening. The setup script for UvA-managed Windows machines works and goes out to the first cohort now, so broken machines show up weeks before the course. Licensing is still open and has a deadline of early next week. The evening 1 introduction will open with motivation and a demo before any folder structure or agent files. Lab 4 will have participants talk to the models directly, not through paper-to-agent wrappers.

Decisions

- **One repository for everything.** Slides, the setup script, transcripts and lab material live in one GitHub repo. Participants get a single URL. The README opens with a step-by-step quick start.
- **Setup runs before the course.** Participants run the PowerShell script from the repo at home. Nothing is asked of them during the run. Machines where it fails get sorted out in advance, and pairs are the fallback on the night.
- **A standard permission file in every lab repo.** The setup script will place it in each lab folder so the agent behaves the same for everyone. It covers pushing to shared repos and handling personal data.
- **Direct model use for replication.** Paper-to-agent tools added little in Joaquim's test and are out of Lab 4.
- **Evening 1 order.** Why the course exists, references and a small demo first. Folder structure and agent files after that.
- **Editorial and faculty AI policy goes in Lab 2** (day 2), alongside the ethics material.

Action items

Owner	Action	When
Joaquim	Settle with Victor whether a collective Anthropic license is possible, collective or individual	Early next week (Monday)
Joaquim	Book the Lab 4 briefing with Yenan and Marleen. Put the faculty AI policy slides and document in the repo	Next week
Joaquim	Try to schedule a second run of the course this year (about 70 people signed up, split into two groups)	Open
Gjorgji	Create the clean repo and send invites. Add the permission file to the setup script	Today
Gjorgji	Share the evening 1 slides with the team so Lab content does not overlap. Finish the slides	Before the next meeting
Gjorgji	Meet Marleen at the university and go through her emails	Next week
Marc (to confirm)	Send the setup script to the first-round participants and tell them to contact Gjorgji if it fails	Now

Owner	Action	When
Marc (to confirm)	Call a follow-up meeting to finalise the evenings	In about 10 days

Cohort and logistics

- About 70 people signed up. They are split into two groups, and the first group is capped at 35. The team expects dropout to leave 25 to 28 active participants.
- Participants get 10 credit hours for the first run. A second run may get fewer, still under discussion. The transcript is garbled here, *to confirm*.
- Gjorgji tested Claude Code, Python and GitHub on a UvA-managed machine and everything installs. Settings can change, so the pair fallback stays.
- Joaquim expects the script to work on about 80 percent of machines. Sending it out now is how the team finds the other 20.
- Participants should not receive many separate emails. They get the repo URL, and the README tells them what to do.

Licensing

The problem Joaquim raised: the course teaches privacy and compliance, then asks participants to work on private licenses with no agreement covering their data. Options discussed:

- **Collective license through the university.** Preferred, because a contract with a legal entity covers participants. Joaquim takes this to Victor.
- **Anthropic organisation account.** Gjorgji pointed out that one can be created without a company or VAT number, with seats for each person. Every user pays API prices, so each would get a quota, for example 200 euro. It does not by itself give a legal data agreement.
- **A contract through a private company.** Mentioned and set aside, since it is not the right legal route either.
- **Plan B.** The committee already allows participants to declare their own costs. If no collective deal exists by early next week, everyone buys individually, and the course says plainly that this is not the way to do it normally.

Evenings

Evening 1: introduction and project setup

- Marc opens with the purpose. The course does not teach end-to-end automated research. It shows what the tools can do, where they are weak, and how to use them responsibly. Nobody should leave thinking they learned to automate paper writing.
- Marleen warned that starting with agent files and folder setup overwhelms people. The team agreed to lead with motivation, research references and a small demo.
- Gjorgji then covers project organisation: a project directory with a logbook and agent instruction files, backed up on GitHub, so that you can answer an editor who asks what exactly you did. This doubles as a first example of a skill, and skills can be set to run without being invoked.
- Permissions. Claude loads a permission file before it starts. Participants should not approve every action by hand. Auto mode (Shift+Tab twice) leaves confirmation for serious steps such as overwriting or deleting files.

Lab 1: article from a paper (Marleen and her partner at FD)

- Her partner is a podcast maker, not a journalist, and does not want a decorative role. A longer slot or a different guest may be needed. The team can also run it without an outside guest.

- Idea for the lab: show levels of AI use on one simple task. First a plain prompt, then an agent that questions you about your core message and audience, then a skill file with your writing style and links to earlier articles. The message is that you choose the level of complexity, and you do not need the highest.
- Gjorgji demoed his *grilling* skill (the Grill Me variant). It asks questions and gives a recommended answer for each until the idea is sharp. It is public in his skills repo of about 40 skills.
- A cleaner skill to remove AI writing patterns also goes in. The statistics of these patterns shift over time, so it needs upkeep.
- Homework: bring a paper you want to turn into an article. The team assumes few will do it, so setup must not depend on it.

Lab 2: privacy, ethics and policy

- The faculty AI policy gets short attention. Most points boil down to using AI sensibly, and the main one is to verify what it gives you.
- Two levels: make people aware that the rules exist and where to find them, then show a concrete application in an exercise. The team wants both.
- Joaquim's short editorial-policy exercise belongs here as optional material.

Lab 3: library tool

- Two goals: think about the future of universities and the participant's field in the age of AI, and write about it using deep research plus a tool the participant builds to search the library.
- Participants write something themselves, use the tools, and compare the results. Verification is the main point.
- API keys are needed. OpenAlex gives one easily, and there may be an MCP server for it, *to confirm*.

Lab 4 and evening 3: paper replication

- Joaquim replicated a paper with both Codex and Claude, and it went smoothly. Paper-to-agent took heavy setup and added no value. Direct interaction with Claude was enough.
- A commercial review tool popular with economists is expensive and gave only superficial comments on a paper just submitted to AER. The team plans to talk about this without recommending it.
- Order for the evening: replication exercise first, then a short showcase. A Lean demo was floated but nobody is committed, and one of the leads is not familiar with it.
- Homework idea: propose an important paper from your own field, start a branch, write prompts, and report whether the result is useful.
- Joaquim, Yenan and Marleen will hold a short briefing next week to organise the evening.

Open questions

- Does the collective license work, and if not, how do we frame the individual route?
- Credit hours: 10 for the first run and what for the second? To confirm.
- Who runs the Lean demo, if anyone?
- Does Marleen's partner take part, and is five minutes enough for the role?
- Course dates and room, still unconfirmed from the 15 September meeting.
- When does the full material go to participants? Consensus is only the setup script goes out early.

Speaker labels in the transcript are unreliable, and two labels each seem to cover more than one person. The auto-summary also mangles names (Georgie, Joachim, Yuki, Mark, Marlene). Attributions here follow context, so check them before sharing.